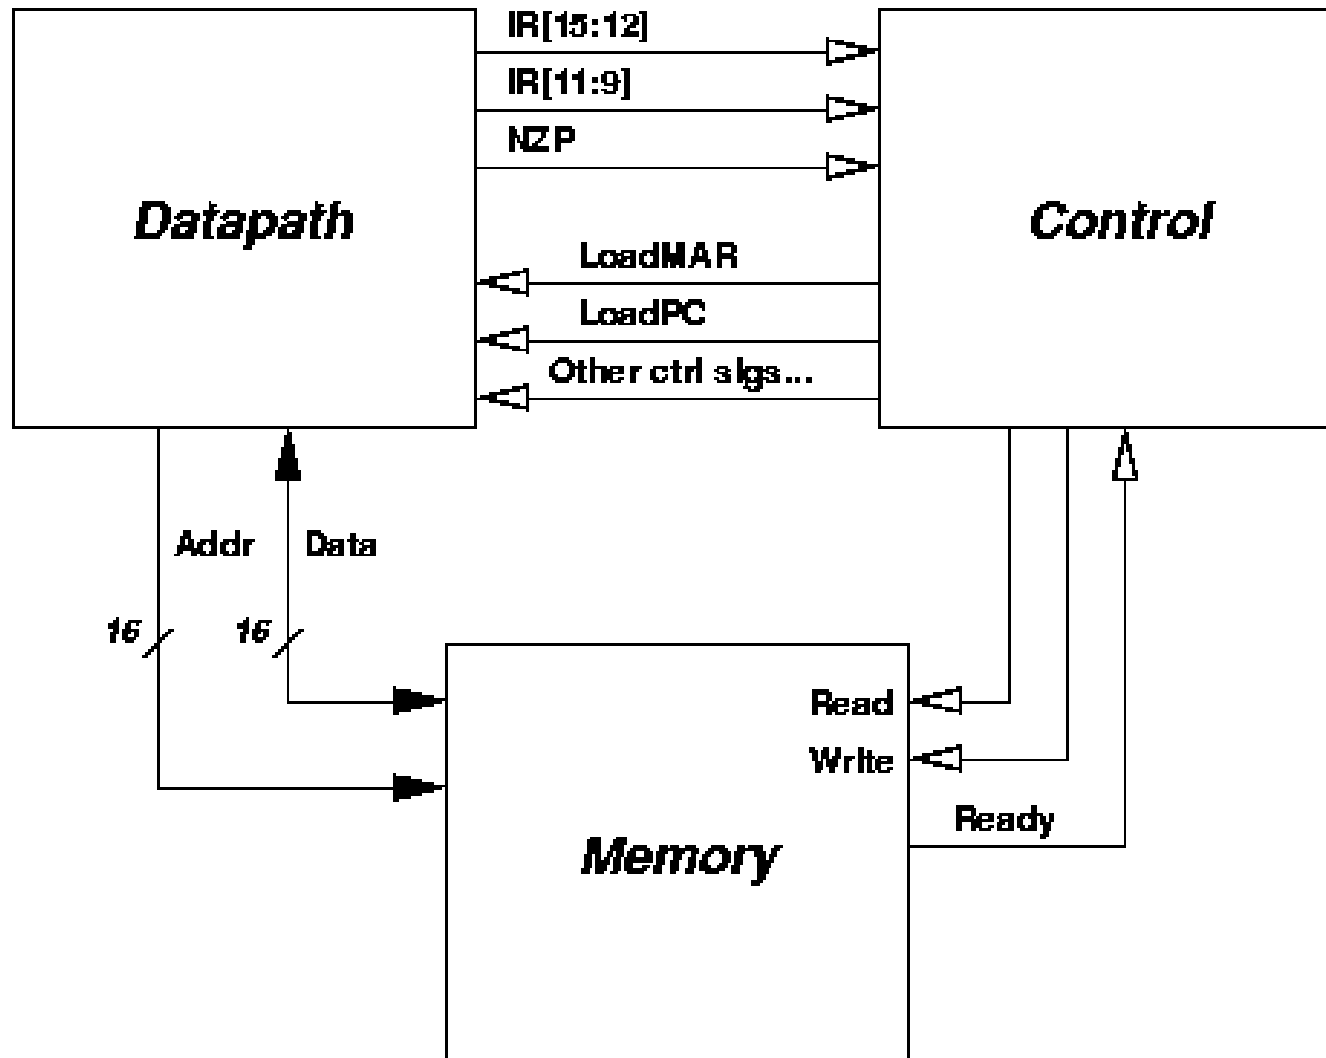



ECE 411  
Fall 2015

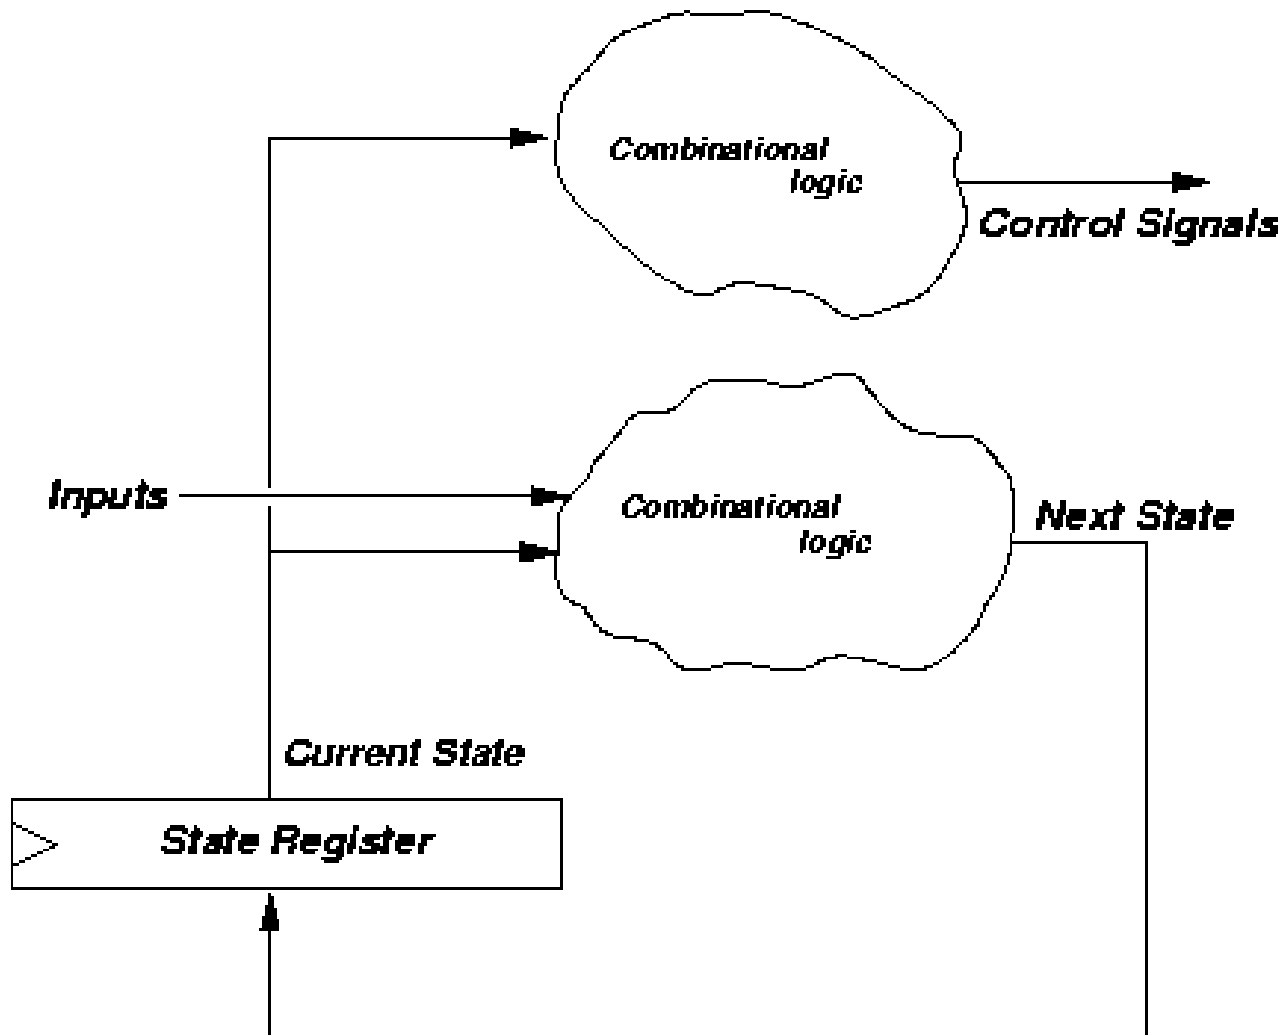
## Lecture 2

Control Unit Design; ISA Design;  
Performance Considerations

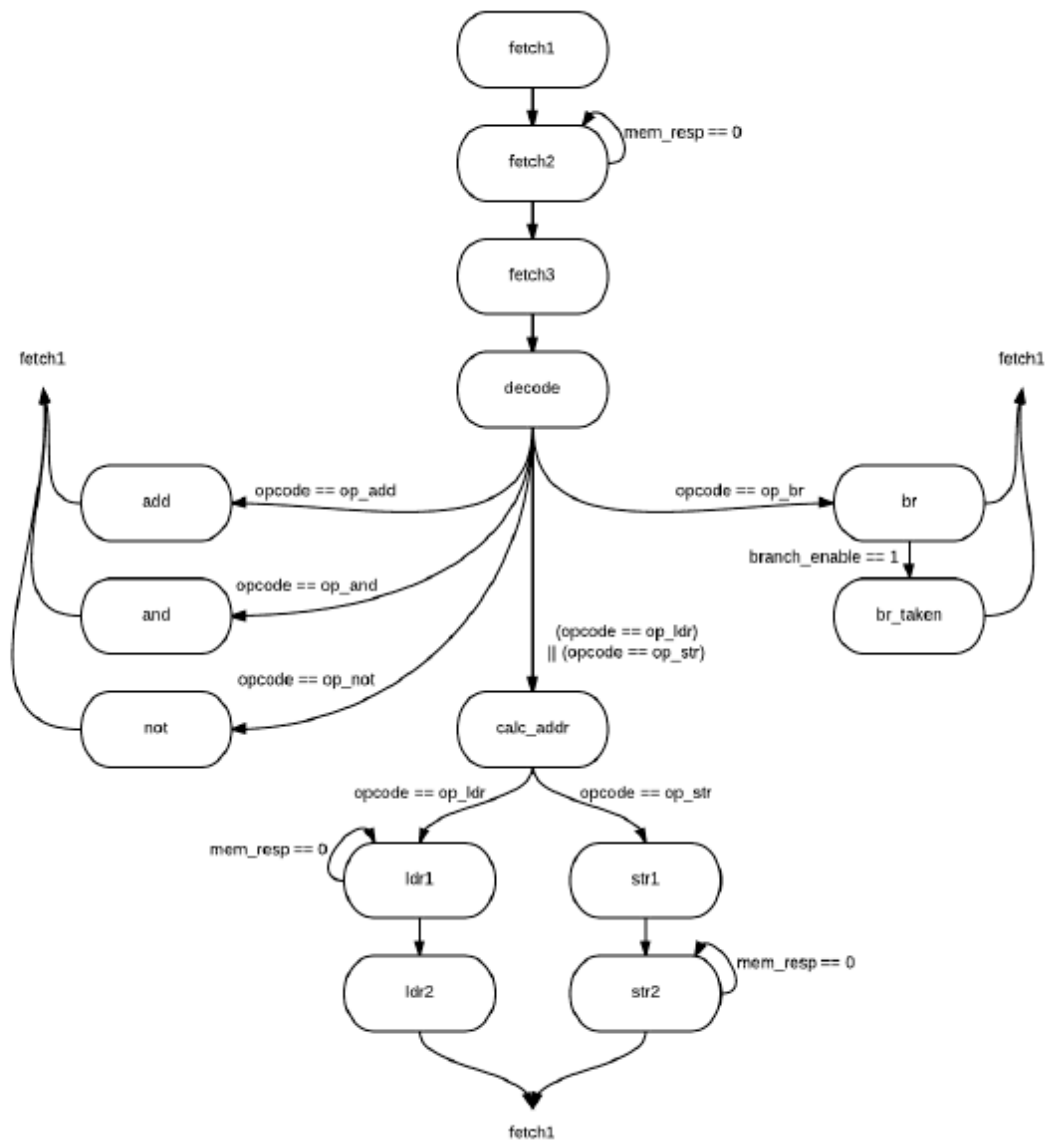
# Control Unit – Control Signals




# Control Signals and Next State




# State Diagram





```
always_comb
begin : state_actions
/* Default assignments */
load_pc = 1'b0;
load_ir = 1'b0;
load_regfile = 1'b0;
aluop = alu_add;
/* ... */
case(state)
  fetch1: begin
    /* MAR <= PC */
    marmux_sel = 1;
    load_mar = 1;
    /* PC <= PC + 2 */
    pcmux_sel = 0;
    load_pc = 1;
  end
end
```

```
  fetch2: begin
    /* Read memory */
    mem_read = 1;
    mdrmux_sel = 1;
    load_mdr = 1;
  end
  fetch3: begin
    /* Load IR */
    load_ir = 1;
  end
  decode: /* Do nothing */;
  s_add: begin
    /* DR <= SRA + SRB */
    aluop = alu_add;
    load_regfile = 1;
    regfilemux_sel = 0;
    load_cc = 1;
  end
end
/*...*/
```



# Next State Logic in RTL Code

```
always_comb
begin : next_state_logic
    next_state = state;
```


```
case(state)
    fetch1: next_state = fetch2;
    fetch2: if (mem_resp) next_state = fetch3;
    fetch3: next_state = decode;
```

```
decode: begin
    case (dp.opcode)
        op_add: next_state = s_add;
        op_and: next_state = s_and;
        ...
    endcase
end
```

```
s_br: begin
    if (dp.nzp_match)
        next_state = s_br_taken;
    else
        next_state = fetch1;
    end

    ...
    default: next_state = fetch1;

endcase
end
```



# Key ISA decisions

---

Instruction format – Length? Variable? Fixed?  
Fields?


How many registers?

Where do instruction operands reside?

- e.g., can you add contents of memory to a register?

Operands

- how many? how big?
- how are memory addresses computed?



# Instruction Length

Variable:

**x86 – Instructions vary from 1 to 17 Bytes long**


**VAX – from 1 to 54 Bytes**

Fixed:

**MIPS, PowerPC, and most other RISC's:**


**all instruction are 4 Bytes long**





# Instruction Length


- Variable-length instructions (x86, VAX):
  - require multi-step, complex fetch and decode.
  - + allow smaller binary programs that require less disk storage, less DRAM at runtime, less memory, bandwidth and better cache efficiency
- Fixed-length instructions (RISC's)
  - + allow easy fetch and decode.
  - + simplify pipelining and parallelism.
  - result in larger binary programs that require more disk storage, more DAM at runtime, more memory bandwidth and lower cache efficiency



# ARM Case Study

- ARM (Advanced RISC Machine)
  - Started with fixed, 32-bit instruction length
  - Added Thumb instructions
    - A subset of the 32-bit instructions
    - All encoded in 16 bits
    - All translated into equivalent 32-bit instructions within the processor pipeline at runtime
    - Can access only 8 general purpose registers
- Motivated by many resource constrained embedded applications that require less disk storage, less DRAM at runtime, less memory, bandwidth and better cache efficiency

# How many registers?



- Most computers have a small set of registers
  - Memory to hold values that will be used soon
  - A typical instruction use 2 or 3 register values
- Advantages of a small number of registers:
  - It requires fewer instruction bits to specify which one.
  - Less hardware
  - Faster access (shorter wires, fewer gates)
  - Faster context switch (when all registers need saving)
- Advantages of a larger number:
  - Fewer loads and stores needed
  - Easier to express several operations in parallel

In 411, “load” means moving data from memory to register, “store” is reverse





# Where do operands reside (*when the ALU needs them*)?

## Stack machine:

“Push” loads memory into 1<sup>st</sup> register (“top of stack”), moves other regs down

“Pop” does the reverse.

“Add” combines contents of first two registers, moves rest up.

## Accumulator machine:

Only 1 register (called the “accumulator”)

Instruction include “store” and “acc  $\leftarrow$  acc + mem”


## Register-Memory machine :

Arithmetic instructions can use data in registers and/or memory

## Load-Store Machine (aka Register-Register Machine):

Arithmetic instructions can only use data in registers.

# Comparing the ISA classes



Code sequence for  $C = A + B$

## Stack

Push A  
Push B  
Add  
Pop C

Java VMs

## Accumulator

Load A  
Add B  
Store C

DSPs

## Register-Memory

Add C, A, B


VAX, x86 partially

## Load-Store

Load R1, A  
Load R2, B  
Add R3, R1, R2  
Store C, R3



# Comparing the ISA classes



Stack

Accumulator

Reg-Mem

Load-store

$$A = X * Y + X * Z$$

Accumulator

R1

R2

R3


Stack

Memory

A	?
X	12
Y	3
B	4
C	5
temp	?



# Load-store architectures



can do:

`add r1=r2+r3`

`load r3, M(address)`

`store r1, M(address)`


⇒ forces heavy dependence on registers, which works for today's CPUs

can't do:

`add r1=r2+M(address)`

- more instructions
- + fast implementation (e.g., easy pipelining)
- + easier to keep instruction lengths fixed





# Instruction formats *-what does each bit mean?*

Machine needs to determine quickly,


- “This is a 6-byte instruction”
- “Bits 7-11 specify a register”
- ...
- Serial decoding bad

Having many different instruction formats...

- complicates decoding
- uses instruction bits (to specify the format)


**What would be a good thing about having many different instruction formats?**







# LC-3b Instruction Formats

- ADD, AND (without Immediate)




- ADD, AND (with Immediate), NOT






# LC-3b Instruction Formats


- ADD, AND (without Immediate)



- ADD, AND (with Immediate), NOT




# MIPS Instruction Formats



	6 bits	5 bits	5 bits	5 bits	5 bits	6 bits
r format	OP	rs	rt	rd	sa	funct
i format	OP	rs	rt	immediate		
j format	OP	target				

- for instance, "add r1, r2, r3" has
  - OP=0, rs=2, rt=3, rd=1, sa=0 (shift amount), funct=32
  - 000000 00010 00011 00001 00000 100000
- opcode (OP) tells the machine which format

# MIPS ISA Tradeoffs




6 bits	5 bits	5 bits	5 bits	5 bits	6 bits
OP	rs	rt	rd	sa	funct
OP	rs	rt	immediate		
OP	target				

What if?

- 64 registers
- 20-bit immediates
- 4 operand instruction (e.g.  $Y = AX + B$ )


Think about how sparsely the bits are used

# Conditional branch




- How do you specify the **destination** of a branch/jump?
  - Theoretically, the destination is a full address
    - 16 bits for LC3b
    - 32 bits for MIPS

# Conditional branch




- How do you specify the destination of a branch/jump?
- studies show that almost all conditional branches go **short distances** from the current program counter (loops, if-then-else).

# Conditional branch



- How do you specify the destination of a branch/jump?
- studies show that almost all conditional branches go short distances from the current program counter (loops, if-then-else).
  - we can specify a relative address in much fewer bits than an absolute address
  - e.g., `beq $1, $2, 100`  $\Rightarrow$  if ( $\$1 == \$2$ )  $PC = PC + 100 * 4$


# Conditional branch



- How do you specify the destination of a branch/jump?
- studies show that almost all conditional branches go short distances from the current program counter (loops, if-then-else).
  - we can specify a relative address in much fewer bits than an absolute address
  - e.g., `beq $1, $2, 100`  $\Rightarrow$  if ( $\$1 == \$2$ )  $PC = PC + 100 * 4$
- How do we specify the **condition** of the branch?




# MIPS conditional branches




- beq, bne     *beq r1, r2, addr => if (r1 == r2) goto addr*
- slt \$1, \$2, \$3   => if (\$2 < \$3) \$1 = 1; else \$1 = 0
- these, combined with \$0, can implement all fundamental branch conditions

Always, never, !=, =, >, <=, >=, <, >(unsigned), <= (unsigned), ...

```
if (i<j)
    w = w+1;
else
    w = 5;
```



```
slt $temp, $i, $j
beq $temp, $0, L1
add $w, $w, #1
beq $0, $0, L2
L1: add $w, $0, #5
L2:
```



# Jumps

- need to be able to jump to an absolute address sometime
- need to be able to do procedure calls and returns
- `jump -- j 10000 => PC = 10000`
- `jump and link -- jal 100000 => $31 = PC + 4; PC = 10000`
  - used for procedure calls

OP	target
----	--------


- `jump register -- jr $31 => PC = $31`
  - used for returns, but can be useful for lots of other things.

A detailed, colorful image of a microchip or integrated circuit, showing various components and patterns in red, yellow, green, and blue.

# Computer Performance


---

# Computer Performance: TIME, TIME, TIME



- Response Time (latency)
    - How long does it take for my job to run?
    - How long does it take to execute a job?
    - How long must I wait for the database query?
  - Throughput
    - How many jobs can the machine run at once?
    - What is the average execution rate?
    - How much work is getting done?
- *If we upgrade a machine with a new processor what do we increase?*
- If we add a new machine to the lab what do we increase?*

# Aspects of CPU Performance



$$\text{CPU time} = \frac{\text{Seconds}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Cycles}}{\text{Instruction}} \times \frac{\text{Seconds}}{\text{Cycle}}$$

	Inst Count	CPI	Clock Rate
Program	X		
Compiler	X	(X)	
Inst. Set.	X	X	
Organization	X	X	X
Technology			X



# P&H Definition of Performance

---

- For some program running on machine X,

$$\text{Performance}_X = 1 / \text{Execution time}_X$$


- "X is n times faster than Y"

$$\text{Performance}_X / \text{Performance}_Y = n$$

- Problem:

- machine A runs a program in 20 seconds
- machine B runs the same program in 25 seconds

# How to Improve Performance




$$\frac{\text{seconds}}{\text{program}} = \frac{\text{cycles}}{\text{program}} \times \frac{\text{seconds}}{\text{cycle}}$$

So, to improve performance (everything else being equal) you can either

\_\_\_\_\_ the # of required cycles for a program, or

\_\_\_\_\_ the clock cycle time or, said another way,

\_\_\_\_\_ the clock rate.




# Now that we understand cycles

- A given program will require
  - some number of instructions (machine instructions)
  - some number of cycles
  - some number of seconds
- We have a vocabulary that relates these quantities:
  - cycle time (seconds per cycle)
  - clock rate (cycles per second)
  - CPI (cycles per instruction)
    - a floating point intensive application might have a higher CPI*
  - MIPS (millions of instructions per second)
    - this would be higher for a program using simple instructions*



# CPI Example



- Suppose we have two implementations of the same instruction set architecture (ISA).


For some program,

Machine A has a clock cycle time of 10 ns. and a CPI of 2.0

Machine B has a clock cycle time of 20 ns. and a CPI of 1.2

What machine is faster for this program, and by how much?

- *If two machines have the same ISA which of our quantities (e.g., clock rate, CPI, execution time, # of instructions, MIPS) will likely be identical during a comparison?*



# # of Instructions Example


- A compiler designer is trying to decide between two code sequences for a particular machine. Based on the hardware implementation, there are three different classes of instructions: Class A, Class B, and Class C, and they require one, two, and three cycles (respectively).

The first code sequence has 5 instructions: 2 of A, 1 of B, and 2 of C

The second sequence has 6 instructions: 4 of A, 1 of B, and 1 of C.

Which sequence will be faster? How much?

What is the CPI for each sequence?




# MIPS example

- Two different compilers are being tested for a 100 MHz. machine with three different classes of instructions: Class A, Class B, and Class C, which require one, two, and three cycles (respectively). Both compilers are used to produce code for a large piece of software.

The first compiler's code uses 5 million Class A instructions, 1 million Class B instructions, and 1 million Class C instructions.

The second compiler's code uses 10 million Class A instructions, 1 million Class B instructions, and 1 million Class C instructions.

- Which sequence will be faster according to MIPS?
- Which sequence will be faster according to execution time?




# Which is faster in the same CPU?

load R1, addr1	→	load R1, addr1
store R1, addr2	→	add R0, R2 -> R3
add R0, R2 -> R3	→	add R0, R6 -> R7
subtract R4, R3 -> R5	→	store R1, addr2
add R0, R6 -> R7	→	subtract R4, R3 -> R5
store R7, addr3	→	store R7, addr3

**Twice as fast on some  
machines and same on others**

**CPI can vary due to interactions among instructions!**



# **READ CHAPTER 2 COMPLETE MP0**

