


ECE 411
Fall 2015

Lecture 3

ISA Design;
Performance Metrics



Key ISA decisions

Instruction format – Length? Variable? Fixed?
Fields?


How many registers?

Where do instruction operands reside?

- e.g., can you add contents of memory to a register?

Operands

- how many? how big?
- how are memory addresses computed?



Where do operands reside (*when the ALU needs them*)?

Stack machine:

“Push” loads memory into 1st register (“top of stack”), moves other regs down

“Pop” does the reverse.

“Add” combines contents of first two registers, moves rest up.

Accumulator machine:

Only 1 register (called the “accumulator”)

Instruction include “store” and “acc \leftarrow acc + mem”


Register-Memory machine :

Arithmetic instructions can use data in registers and/or memory

Load-Store Machine (aka Register-Register Machine):

Arithmetic instructions can only use data in registers.

Comparing the ISA classes



Code sequence for $C = A + B$

Stack

Push A
Push B
Add
Pop C

Java VMs

Accumulator

Load A
Add B
Store C

DSPs

Register-Memory


Add C, A, B

VAX, x86 partially

Load-Store

Load R1, A
Load R2, B
Add R3, R1, R2
Store C, R3

Comparing the ISA classes



Stack

Accumulator

Reg-Mem

Load-store

$$A = X * Y + X * Z$$

Accumulator

R1

R2


R3

Stack

Memory

A	?
X	12
Y	3
B	4
C	5
temp	?





Instruction formats *-what does each bit mean?*


Machine needs to determine quickly,

- “This is a 6-byte instruction”
- “Bits 7-11 specify a register”
- ...
- Serial decoding bad

Having many different instruction formats...


- complicates decoding
- uses instruction bits (to specify the format)

What would be a good thing about having many different instruction formats?





LC-3b Instruction Formats

- ADD, AND (without Immediate)




- ADD, AND (with Immediate), NOT






LC-3b Instruction Formats


- ADD, AND (without Immediate)



- ADD, AND (with Immediate), NOT




MIPS Instruction Formats



	6 bits	5 bits	5 bits	5 bits	5 bits	6 bits
r format	OP	rs	rt	rd	sa	funct
i format	OP	rs	rt	immediate		
j format	OP	target				

- for instance, "add r1, r2, r3" has
 - OP=0, rs=2, rt=3, rd=1, sa=0 (shift amount), funct=32
 - 000000 00010 00011 00001 00000 100000
- opcode (OP) tells the machine which format

MIPS ISA Tradeoffs




6 bits	5 bits	5 bits	5 bits	5 bits	6 bits
OP	rs	rt	rd	sa	funct
OP	rs	rt	immediate		
OP	target				

What if?

- 64 registers
- 20-bit immediates
- 4 operand instruction (e.g. $Y = AX + B$)


Think about how sparsely the bits are used

Conditional branch




- How do you specify the **destination** of a branch/jump?
 - Theoretically, the destination is a full address
 - 16 bits for LC3b
 - 32 bits for MIPS

Conditional branch




- How do you specify the destination of a branch/jump?
- studies show that almost all conditional branches go **short distances** from the current program counter (loops, if-then-else).

Conditional branch




- How do you specify the destination of a branch/jump?
- studies show that almost all conditional branches go short distances from the current program counter (loops, if-then-else).
 - we can specify a relative address in much fewer bits than an absolute address
 - e.g., `beq $1, $2, 100` \Rightarrow if ($\$1 == \2) $PC = PC + 100 * 4$

Conditional branch



- How do you specify the destination of a branch/jump?
- studies show that almost all conditional branches go short distances from the current program counter (loops, if-then-else).
 - we can specify a relative address in much fewer bits than an absolute address
 - e.g., `beq $1, $2, 100` \Rightarrow if ($\$1 == \2) $PC = PC + 100 * 4$
- How do we specify the **condition** of the branch?


MIPS conditional branches




- beq, bne *beq r1, r2, addr => if (r1 == r2) goto addr*
- slt \$1, \$2, \$3 => if (\$2 < \$3) \$1 = 1; else \$1 = 0
- these, combined with \$0, can implement all fundamental branch conditions

Always, never, !=, ==, >, <=, >=, <, >(unsigned), <= (unsigned), ...

```
if (i<j)
    w = w+1;
else
    w = 5;
```



```
slt $temp, $i, $j
beq $temp, $0, L1
add $w, $w, #1
beq $0, $0, L2
L1: add $w, $0, #5
L2:
```




Jumps


- need to be able to jump to an absolute address sometimes
- need to be able to do procedure calls and returns
- `jump -- j 10000 => PC = 10000`
- `jump and link -- jal 100000 => $31 = PC + 4; PC = 10000`
 - used for procedure calls

OP	target
----	--------

- `jump register -- jr $31 => PC = $31`
 - used for returns, but can be useful for lots of other things.



Computer Performance



Computer Performance: TIME, TIME, TIME


- Response Time (latency)
 - How long does it take to execute a job?
 - How long must I wait for the database query?
- Throughput
 - How many jobs can the machine complete in a minute?
 - What is the average execution rate?
 - How much work is getting done?
- *If we upgrade a machine with a new processor what do we improve?*
- *If we add a new machine to the lab what do we improve?*

A Throughput Oriented Approach

- Each car still experience the same latency at the toll
- Many cars are served to improve the number of cars processed per time unit (sec/min/hr)




Case Study



- How can we improve the throughput of I-57 between Chicago and Champaign?
 - Increase the number of lanes
- How can we improve the latency of driving from Chicago to Champaign on I-57?
 - Increase the driving speed

Aspects of CPU Performance



$$\text{CPU time} = \frac{\text{Seconds}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Cycles}}{\text{Instruction}} \times \frac{\text{Seconds}}{\text{Cycle}}$$

	Inst Count	CPI	Clock Rate
Program	X		
Compiler	X	X	
Inst. Set.	X	X	
Organization	X	X	X
Technology			X



P&H Definition of Performance

- For some program running on machine X,

$$\text{Performance}_X = 1 / \text{Execution time}_X$$


- "X is n times faster than Y"

$$\text{Performance}_X / \text{Performance}_Y = n$$

- Problem:

- machine A runs a program in 20 seconds
- machine B runs the same program in 25 seconds
- A is 1.25 times faster than B?

How to Improve Performance




$$\frac{\text{seconds}}{\text{program}} = \frac{\text{cycles}}{\text{program}} \times \frac{\text{seconds}}{\text{cycle}}$$

So, to improve performance (everything else being equal) you can either

_____ the # of required cycles for a program, or


_____ the clock cycle time or, said another way,

_____ the clock rate.



Performance Related Metrics


- An execution of a given program will require
 - some number of instructions (machine instructions)
 - some number of cycles
 - some number of seconds
- We have a vocabulary of metrics that relate these quantities:
 - cycle time (seconds per cycle)
 - clock rate (cycles per second)
 - CPI (cycles per instruction)
 - a floating point intensive application might have a higher CPI*
 - MIPS (millions of instructions per second)
 - this would be higher for a program using simple instructions*



Instruction Set Metrics Example

- *If two machines have the same ISA which of our quantities (e.g., clock rate, CPI, execution time, # of instructions, MIPS) will likely be identical during a comparison?*

CPI Example



- Suppose we have two implementations of the same instruction set architecture (ISA).

For some program,

Machine A has a clock cycle time of 10 ns. and a CPI of 2.0

Machine B has a clock cycle time of 20 ns. and a CPI of 1.2

What machine is faster for this program, and by how much?




Cycles and CPI Example

- A compiler designer is trying to decide between two code sequences for a particular machine. Based on the hardware implementation, there are three different classes of instructions: Class A, Class B, and Class C, and they require one, two, and three cycles (respectively).

The first code sequence has 5 instructions: 2 of A, 1 of B, and 2 of C

The second sequence has 6 instructions: 4 of A, 1 of B, and 1 of C.

- Which sequence will be faster? How much?
 - First: $2*1 + 1*2 + 2*3 = 10$ cycles
 - Second: $4*1 + 1*2 + 1*3 = 9$ cycles
 - Second is faster than first by 1 cycle or about 11%
- What is the average CPI for each sequence?
 - First: $10/5 = 2$
 - Second: $9/6 = 1.5$




MIPS example

- Two different compilers are being tested for a 100 MHz. machine with three different classes of instructions: Class A, Class B, and Class C, which require one, two, and three cycles (respectively). Both compilers are used to produce code for a large piece of software.

The first compiler's code uses 5 million Class A instructions, 1 million Class B instructions, and 1 million Class C instructions.

The second compiler's code uses 10 million Class A instructions, 1 million Class B instructions, and 1 million Class C instructions.

- Which sequence will have higher MIPS rating?
 - First: $(5M + 1M + 1M) / ((5M \cdot 1 + 1M \cdot 2 + 1M \cdot 3) / 100M) = 70 \text{ MIPS}$
 - Second: $(10M + 1M + 1M) / ((10M \cdot 1 + 1M \cdot 2 + 1M \cdot 3) / 100M) = 80 \text{ MIPS}$




MIPS example (cont.)

- Two different compilers are being tested for a 100 MHz. machine with three different classes of instructions: Class A, Class B, and Class C, which require one, two, and three cycles (respectively). Both compilers are used to produce code for a large piece of software.

The first compiler's code uses 5 million Class A instructions, 1 million Class B instructions, and 1 million Class C instructions.


The second compiler's code uses 10 million Class A instructions, 1 million Class B instructions, and 1 million Class C instructions.

- Which sequence will be faster according to execution time?
 - First: $5M \times 1 + 1M \times 2 + 1M \times 3 = 10M$ cycles
 - Second: $10M \times 1 + 1M \times 2 + 1M \times 3 = 15M$ cycles



MP1 - Important!

- All remaining MP due dates include an automatic 2-day extension.
 - Treat Fridays as the real deadline!
 - Extended due dates are on Sundays. Use the two days as contingency.
- Read the document carefully before you start!
- Review relevant parts of the document as you progress!
- Follow instructions and steps!



READ 1.6 AND CHAPTER 2
START MP1

